

Lab8: Locks

2351882 王小萌

Tongji University,2025 Summer

代码仓库链接: <https://github.com/wmxmw/Tongji-University-xv6-labs-2021.git>

1. Memory allocator

1.1 实验目的

1.2 实验步骤

1.3 实验中遇到的问题和解决办法

1.4 实验心得

2.Buffer Cache

2.1 实验目的

2.2 实验步骤

2.3 实验中遇到的问题和解决办法

2.4 实验心得

3 实验检验得分

1 Memory allocator

1.1 实验目的

为了减少多核系统中的锁竞争并提升性能，可以重构内存分配器的设计。具体做法是为每个 CPU 分配一个 独立的自由列表 (free list)，并且每个自由列表都有 专属的锁。这样，不同 CPU 上的内存分配和释放 操作可以并行执行，显著减少 锁争用。同时，当某个 CPU 的自由列表耗尽时，它应能够从其他 CPU 的自由 列表中获取部分内存。

1.2 实验步骤

1. 在 xv6 中运行 kalloc_test, 输出如下:

```

hart 2 starting
hart 1 starting
init: starting sh
$ kallocatest
exec kallocatest failed
$ kallocatest
start test1
test1 results:
--- lock kmem/bcache stats
lock: kmem: #test-and-set 98472 #acquire() 433070
lock: bcache: #test-and-set 0 #acquire() 1376
--- top 5 contended locks:
lock: proc: #test-and-set 163489 #acquire() 4215604
lock: proc: #test-and-set 162860 #acquire() 4215373
lock: proc: #test-and-set 160260 #acquire() 4215518
lock: proc: #test-and-set 151686 #acquire() 4215574
lock: proc: #test-and-set 143763 #acquire() 4215605
tot= 98472
test1 FAIL
start test2
total free number of pages: 32499 (out of 32768)
.....
test2 OK
$

```

可以看出，test1 测试未通过。并且针对 kmem 锁的争用情况：acquire() 函数被调用了 433070 次，自旋尝试获取锁的次数为 98472 次。锁 proc 也有大量的 #test-and-set 和 #acquire() 操作。

2. 观察到 NCPU 在 param.h 中声明，值为 8。

```

kernel > C param.h
1  #define NPROC      64 // maximum number of processes
2  #define NCPU ..... 8 // maximum number of CPUs
3  #define NOFILE     16 // open files per process
4  #define NFILE      100 // open files per system

```

3. 将 kalloc.c 中 kmem 修改成数组形式结构体，每个 CPU 分配一个 keme 锁：

```

21  struct {
22      struct spinlock lock;
23      struct run *freelist;
24  } kmem[NCPU];
25

```

4. 修改 kalloc.c 中的 kinit() 函数： 初始化每个 CPU 的空闲链表和锁

```

26 void
27 kinit()
28 {
29     int i;
30     for (i = 0; i < NCPU; ++i) {
31         //snprintf(kmem[i].lockname, 8, "kmem_%d", i); //the name of the lo
32         initlock(&kmem[i].lock, "kmem");
33     }
34     // initlock(&kmem.lock, "kmem"); // lab8-1
35     freerange(end, (void*)PHYSTOP);
36 }

```

5. 修改 kalloc.c 中的 kfree() 函数, 确保每个 CPU 的空闲链表正确维护。

```

51 void
52 kfree(void *pa)
53 {
54     struct run *r;
55     int c; // cpuid - lab8-1
56
57     if(((uint64)pa % PGSIZE) != 0 || (char*)pa < end || (uint64)pa >= P
58         panic("kfree");
59
60     // Fill with junk to catch dangling refs.
61     memset(pa, 1, PGSIZE);
62
63     r = (struct run*)pa;
64     // get the current core number - lab8-1
65     push_off();
66     c = cpuid();
67     pop_off();
68     // free the page to the current cpu's freelist - lab8-1
69     acquire(&kmem[c].lock);
70     r->next = kmem[c].freelist;
71     kmem[c].freelist = r;
72     release(&kmem[c].lock);
73 }
74

```

6. 修改 kalloc.c 中的 kalloc() 函数, 分配内存页

```

78 void *
79 kalloc(void)
80 {
81     struct run *r;
82     int c;
83     push_off();
84     c = cpuid();
85     pop_off();
86
87     acquire(&kmem.lock);
88     r = kmem.freelist;
89
90     if(r)
91     {
92         kmem[c].freelist = r->next;
93         release(&kmem[c].lock);
94         if(!r & (r=steal(c))){
95             acquire(&kmem[c].lock);
96             kmem[c].freelist = r->next;
97             release(&kmem[c].lock);
98         }
99         if(r)
100             memset((char*)r, 5, PGSIZE); // fill with junk
101     }
102     return (void*)r;

```

7. 编写 kalloc.c 中的 steal() 函数,从其他 CPU 偷取部分空闲内存页,并添加声明

```

103 struct run *steal(int cpu_id){
104     int i;
105     int c = cpu_id;
106     struct run *fast,*slow,*head;
107     if(cpu_id != cpuid()){
108         panic("steal: cpu_id err");
109     }
110     for(int i=1; i<NCPU; ++i){
111         if(++c == NCPU){
112             c=0;
113         }
114         acquire(&kmem[c].lock);
115         if(kmem[c].freelist){
116             slow = head = kmem[c].freelist;
117             fast = slow->next;
118             while(fast){
119                 fast=fast->next;
120                 if(fast){
121                     slow=slow->next;
122                     fast=fast->next;
123                 }
124             }
125             //the latter half change to current cpu freelist
126             kmem[c].freelist = slow->next;
127             release(&kmem[c].lock);
128             //clear the first half
129             slow->next=0;
130             return head;
131         }
132         release(&kmem[c].lock);
133     }
134     return 0;

```

```

kernel > C kalloc.c
 9  #include "riscv.h"
10  #include "defs.h"
11  struct run *steal(int cpu_id);
12  void freerange(void *pa_start, void *pa_end);

```

8. 运行 xv6，并在命令行中输入 kalloc_test，可以看到 acquire 的循环迭代次数明显减少，表明锁的争用较之前有所减少。

```

xv6 kernel is booting
hart 2 starting
hart 1 starting
init: starting sh
$ kalloc_test
start test1
test1 results:
--- lock kmem/bcache stats
lock: kmem: #test-and-set 0 #acquire() 77629
lock: kmem: #test-and-set 0 #acquire() 182911
lock: kmem: #test-and-set 0 #acquire() 172481
lock: kmem: #test-and-set 0 #acquire() 2
lock: kmem: #test-and-set 0 #acquire() 2
lock: kmem: #test-and-set 0 #acquire() 2
lock: kmem: #test-and-set 0 #acquire() 2
lock: kmem: #test-and-set 0 #acquire() 2
lock: bcache: #test-and-set 0 #acquire() 340
--- top 5 contended locks:
lock: proc: #test-and-set 296689 #acquire() 5668557
lock: proc: #test-and-set 289378 #acquire() 5668529
lock: proc: #test-and-set 235224 #acquire() 5668529
lock: proc: #test-and-set 227207 #acquire() 5668528
lock: proc: #test-and-set 224774 #acquire() 5668529
tot= 0
test1 OK
start test2
total free number of pages: 32499 (out of 32768)
.....
test2 OK
$

```

9. 运行 usertests 测试，确保所有的用户测试都通过

```

test opentest: OK
test writetest: OK
test writebig: OK
test createtest: OK
test openiput: OK
test exitiput: OK
test iput: OK
test mem: OK
test pipel: OK
test killstatus: OK
test preempt: kill... wait... OK
test exitwait: OK
test rmdot: OK
test fourteen: OK
test bigfile: OK
test dirfile: OK
test iref: OK
test forktest: OK
test bigdir: OK
ALL TESTS PASSED
$

```

10. 单项检测得分:

```

vale@Puppyyoc: /xv6-labs-2021$ ./grade-lab-lock kallocetest
make: 'kernel/kernel' is up to date.
== Test running kallocetest == (37.6s)
== Test kallocetest: test1 ==
    kallocetest: test1: OK
== Test kallocetest: test2 ==
    kallocetest: test2: OK
== Test kallocetest: sbrkmuch == kallocetest: sbrkmuch: OK (5.1s)

```

1.3 实验中遇到的问题和解决办法

问题：读取 CPU 的 ID 时没有禁用中断，使得程序出错。。

解决办法：在实现过程中，需要确保每次获取锁和释放锁的顺序是一致的，避免出现循环等待的情况。除此之外，这很有可能是多个核心同时访问共享资源而没有正确的同步机制，也许会导致竞争、数据损坏、不一致性或不可预测的结果。

1.4 实验心得

本次实验通过优化锁的使用，不仅提高了内存分配的性能，还加深了我对操作系统中并发控制的理解。

2 Buffer cache

2.1 实验目的

1. 优化 xv6 操作系统中的缓冲区缓存 (buffer cache)，减少多个进程之间对缓冲区缓存锁的争用，从而提高系统的性能和并发能力。通过设计和实现一种更加高效的缓冲区管理机制，使得不同进程可以更有效地使用和管理缓冲区，减少锁竞争和性能瓶颈
2. 原始的 xv6 中，对于缓存的读写是由单一的锁 bcache.lock 来保护的，导致了如果系统中有多个进程在进行 IO 操作，则等待获取 bcache.lock 的开销就会较大。为了减少加锁的开销，可以将缓存分为几个桶，为每个桶单独设置一个锁，这样如果两个进程访问的缓存块在不同的桶中，则可以同时获得两个锁进行操作，而无需等待加锁。目标是将 bcachetest 中统计值 tot 降到规定值以下。。

2.2 实验步骤

1. 运行 bcachetest 测试程序，观察锁竞争情况和测试结果

```
xv6 kernel is booting

hart 1 starting
hart 2 starting
init: starting sh
$ bcachetest
start test0
test0 results:
--- lock kmem/bcache stats
lock: kmem: #test-and-set 0 #acquire() 32872
lock: kmem: #test-and-set 0 #acquire() 98
lock: kmem: #test-and-set 0 #acquire() 71
lock: kmem: #test-and-set 0 #acquire() 2
lock: kmem: #test-and-set 0 #acquire() 2
lock: kmem: #test-and-set 0 #acquire() 2
lock: kmem: #test-and-set 0 #acquire() 2
lock: kmem: #test-and-set 0 #acquire() 2
lock: bcache: #test-and-set 29754 #acquire() 72492
--- top 5 contended locks:
lock: virtio_disk: #test-and-set 126050 #acquire() 1236
lock: proc: #test-and-set 126014 #acquire() 3075793
lock: proc: #test-and-set 109598 #acquire() 3096368
lock: proc: #test-and-set 107231 #acquire() 3096370
lock: proc: #test-and-set 105411 #acquire() 3096311
tot= 29754
test0: FAIL
start test1
test1 OK
$
```

2. 设计和实现缓冲区管理机制，修改 buf 和 bcache 结构体，为每个缓冲区添加时间戳字段。为 bcache 添加哈希表的锁。

```

kernel > C buf.h
1  struct buf {
2      int valid;    // has data been read from disk?
3      int disk;    // does disk "own" buf?
4      uint dev;
5      uint blockno;
6      struct sleeplock lock;
7      uint refcnt;
8
9      struct buf *prev; // LRU cache list
10     struct buf *next; //hash list
11     uchar data[BSIZE];
12     uint timestamp//the buf last using time -lab8-2
13 };
14
15

```

```

25
26 #define NBUCKET 13
27 #define HASH(blockno)(blockno & NBUCKET)
28
29 struct {
30     struct spinlock lock;
31     struct buf buf[NBUF];
32     int size; //the count of used buf -lab8-2
33     struct buf buckets[NBUCKET];
34     struct spinlock locks[NBUCKET]; //buckets' locks
35     struct spinlock hashlock; //the hash table's lock
36
37     // Linked list of all buffers, through prev/next.
38     // Sorted by how recently the buffer was used.
39     // head.next is most recent, head.prev is least.
40     // struct buf head;
41 } bcache;
42
43 void

```

3. 修改初始化函数。初始化哈希表的锁和缓冲区的时间戳字段

```

43 void
44 binit(void)
45 {
46     int i;
47     struct buf *b;
48     bcache.size=0;
49
50     initlock(&bcache.lock, "bcache");
51     initlock(&bcache.hashlock, "bcache_hash"); //init hashlock
52
53     //init all buckets' lock
54     for(i=0;i<NBUCKET;++i){
55         initlock(&bcache.lock[i], "bcache_bucket");
56     }
57
58     for(b = bcache.buf; b < bcache.buf+NBUF; b++){
59         initsleeplock(&b->lock, "buffer");
60     }
61 }
62
63

```


4. 修改 brelse() 函数： 释放缓冲区时，更新时间戳字段，根据时间戳来选择缓冲区进行重用。

```
125 void
126 brelse(struct buf *b)
127 {
128     int idx;
129     if(!holdingsleep(&b->lock))
130         panic("brelse");
131
132     releasesleep(&b->lock);
133
134     idx =HASH(b->blockno)
135     acquire(&bcache.lock[idx]);
136     b->refcnt--;
137     if (b->refcnt == 0) {
138         // no one is waiting for it.
139         b->timestamp = ticks;
140     }
141
142     release(&bcache.lock);
143 }
```

5. 修改分配函数 bget() 函数，使用哈希表来寻找对应的缓冲区，

```

68 bget(uint dev, uint blockno)
69 {
70     struct buf *b;
71     //lab8-2
72     int idx=HASH(blockno);
73     struct buf *pre,*minb = 0,*minpre;
74     uint mintimestamp;
75     int i;
76
77     //loop up
78     acquire(&bcache.locks[idx]);
79     for(b = bcache.buckets[idx].next; b; b = b->next){
80         if(b->dev == dev && b->blockno == blockno){
81             b->refcnt++;
82             release(&bcache.locks[idx]);
83             acquiresleep(&b->lock);
84             return b;
85         }
86     }
87
88     // Not cached.
89     // check if there is a buf not used
90     acquire(&bcache.lock);
91     if(bcache.size<NBUF){
92         b=&bcache.buf[bcache.size++];
93         release(&bcache.lock);
94         b->dev = dev;
95         b->blockno = blockno;
96         b->valid = 0;
97         b->refcnt = 1;
98         b->next = bcache.buckets[idx].next;
99         bcache.buckets[idx].next=b;

```

6. 运行 xv6, 并运行 bcachetest

结果如下:

```

xv6 kernel is booting

hart 2 starting
hart 1 starting
init: starting sh
$ bcachetest
start test0
test0 results:
--- lock kmem/bcache stats
lock: kmem: #test-and-set 0 #acquire() 32934
lock: kmem: #test-and-set 0 #acquire() 49
lock: kmem: #test-and-set 0 #acquire() 54
lock: kmem: #test-and-set 0 #acquire() 1
lock: kmem: #test-and-set 0 #acquire() 1
lock: kmem: #test-and-set 0 #acquire() 1
lock: kmem: #test-and-set 0 #acquire() 1
lock: kmem: #test-and-set 0 #acquire() 1
lock: bcache: #test-and-set 0 #acquire() 84
lock: bcache_hash: #test-and-set 0 #acquire() 54
lock: bcache_bucket: #test-and-set 0 #acquire() 2112
lock: bcache_bucket: #test-and-set 0 #acquire() 4130
lock: bcache_bucket: #test-and-set 0 #acquire() 2268
lock: bcache_bucket: #test-and-set 0 #acquire() 4266
lock: bcache_bucket: #test-and-set 0 #acquire() 4316
lock: bcache_bucket: #test-and-set 0 #acquire() 6315
lock: bcache_bucket: #test-and-set 0 #acquire() 6648
lock: bcache_bucket: #test-and-set 0 #acquire() 7931
lock: bcache_bucket: #test-and-set 0 #acquire() 6179
lock: bcache_bucket: #test-and-set 0 #acquire() 6187
lock: bcache_bucket: #test-and-set 0 #acquire() 6181
lock: bcache_bucket: #test-and-set 0 #acquire() 4121
lock: bcache_bucket: #test-and-set 0 #acquire() 4122
--- top 5 contended locks:
lock: virtio_disk: #test-and-set 99556 #acquire() 1125
lock: proc: #test-and-set 39215 #acquire() 977840
lock: proc: #test-and-set 37494 #acquire() 977840
lock: proc: #test-and-set 36899 #acquire() 977839
lock: proc: #test-and-set 36860 #acquire() 998365
tot= 0
test0: OK
start test1
test1 OK
$ _

```

7. 单项测试得分:

```

vale@Puppyoo: /xv6-labs-2021$ ./grade-lab-lock bcachetest
make: 'kernel/kernel' is up to date.
== Test running bcachetest == (5.6s)
== Test  bcachetest: test0 ==
  bcachetest: test0: OK
== Test  bcachetest: test1 ==
  bcachetest: test1: OK
vale@Puppyoo: ~/xv6-labs-2021$ _

```

2.3 实验中遇到的问题和解决办法

问题：在修改缓冲区管理代码时，可能会遇到并发情况下的数据一致性问题，如资源分配冲突、竞争条件。


解决办法：增加了锁，根据缓冲区的块号计算哈希索引 v ，获取对应的分桶结构体指针 `bucket` 和分桶的自旋锁，以确保在操作缓冲区引用计数时不会被其他线程干扰。

2.4 实验心得

通过本次实验，我通过优化缓冲区管理方式，学会了如何在实际系统中使用更高效的数据结构和算法来提升性能。

3 实验检验得分

1. 在实验目录下创建 `time.txt`，填写完成实验时间数。

 `time.txt`

2025/8/26 19:47

文本文档

1 KB

2. 创建 `answers-lock` 文件将程序运行结果填入。
3. 在终端中执行 `make grade`,得到 lab8 总分：

```
make[1]: Leaving directory '/home/vale/xv6-labs-2021'
== Test running kalloc test ==
$ make qemu-gdb
(37.3s)
== Test kalloc test: test1 ==
kalloc test: test1: OK
== Test kalloc test: test2 ==
kalloc test: test2: OK
== Test kalloc test: sbrkmuch ==
$ make qemu-gdb
kalloc test: sbrkmuch: OK (4.8s)
== Test running bcachetest ==
$ make qemu-gdb
(4.7s)
== Test bcachetest: test0 ==
bcachetest: test0: OK
== Test bcachetest: test1 ==
bcachetest: test1: OK
== Test usertests ==
$ make qemu-gdb
usertests: OK (71.8s)
== Test time ==
time: OK
Score: 70/70
vale@Puppyoo: ~/xv6-labs-2021$
```

4. 代码提交

```
vale@Puppyyoo:~/xv6-labs-2021$ git add .
vale@Puppyyoo:~/xv6-labs-2021$ git commit -m"lab8"
[lock 9f2aa23] lab8
 9 files changed, 384 insertions(+), 56 deletions(-)
 create mode 100644 answers-lock.txt
 create mode 100644 kernel/bio.c:Zone.Identifier
 create mode 100644 kernel/buf.h:Zone.Identifier
 create mode 100644 kernel/kalloc.c:Zone.Identifier
 create mode 100644 kernel/param.h:Zone.Identifier
 create mode 100644 time.txt
vale@Puppyyoo:~/xv6-labs-2021$ git status
On branch lock
Your branch is ahead of 'origin/lock' by 1 commit.
  (use "git push" to publish your local commits)

nothing to commit, working tree clean
vale@Puppyyoo:~/xv6-labs-2021$
```